

Progressive Evolution of GPU Kernels from MatMul to Recurrent Intelligence

SPAK: Systematic Paradigms for Agent based Kernel engineering

Abstract: This project demonstrates a systematic approach to GPU kernel engineering and deep learning architecture design using Semiformal DSL (Domain Specific Language) as the core medium for semantic communication between AI agents. We developed the **LoopLM** architecture, which utilizes temporal recurrence instead of spatial depth to achieve superior algorithmic generalization. Using a Semiformal DSL-based dual-agent paradigm, we demonstrate that a 1-layer recurrent model can outperform a 12-layer static transformer in out-of-distribution arithmetic tasks while using 12x fewer parameters.

1. The SPAK Methodology: The Dual-Agent Paradigm

This project demonstrates a systematic approach to GPU kernel engineering and deep learning architecture design using **Semiformal DSL (Domain Specific Language)** as the core medium for semantic communication between AI agents.

The Dual-Agent Paradigm

LLM agents operate in two distinct specialized roles, synchronized through the DSL:

- **System Engineer (Architect):** Responsible for high-level design, DSL definition, and defining the "laws of physics" for the model.
- **Kernel Engineer (Implementer):** Responsible for low-level GPU kernel implementation (cuTile/CUDA) and conducting error-free experiments.

2. The SPAK Progressive Kernel Roadmap

Our research evolved through four distinct phases, each building upon hardware insights. All kernels were implemented using the **NVIDIA cuTile Python DSL**, with optimizations guided by the **TileGym reference code** (<https://github.com/NVIDIA/TileGym>).

2.1. MatMul: The Hardware Foundation

Optimized basic operation $C = A \times B$ focusing on **Tiling**, **Shared Memory Swizzling**, and **Pipelining** on RTX5070 Blackwell architectures.

$$C_{ij} = \sum_{k=0}^{K-1} A_{ik} \cdot B_{kj}$$

2.2. FMHA: Fusing for Bandwidth

Minimizing HBM traffic by fusing Softmax and Attention into a single kernel.

$$\text{Attention}(Q, K, V) = \text{Softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V$$

2.3. nanoGPT: Positional Geometric Logic

Integration of **Rotary Position Embeddings (RoPE)** for relative token distances logic.

$$q'_m = R_{\Theta,m} q_m, \quad k'_n = R_{\Theta,n} k_n$$

2.4. LoopLM: Temporal Depth

Replacing spatial layers with a recurrent loop where state h evolves L times:

$$h_{l+1} = \text{Block}(h_l), \quad (\text{where } \text{inject_}x_0 = \text{False})$$

3. LoopLM Experimental Analysis

3.1. Experimental Case

Goal: The primary objective of the LoopLM experiments is to demonstrate that intelligence scales more efficiently through temporal recurrence than spatial depth. By repeating a single shared block, we aim to achieve 'Algorithmic Grokking' on tasks that require sequential logic, such as multi-digit addition, while significantly reducing parameter count.

- **GPT-12L (Static):** A standard Transformer model with 12 spatial layers. Primary baseline for fixed-depth architectures.
- **LoopLM-12 (Dynamic):** A recurrent model using 1 shared layer repeated 12 times. Demonstrates temporal depth efficiency.
- **LoopLM-30 (Deep Thinking):** A recurrent model with 30 loops to test the boundaries of 8-digit addition generalization.
- **LoopLM-128e (Efficient):** Extremely compressed version (128e), proving recurrent logic requires fewer parameters.
- **LoopLM-12 (Test-Time 24):** A robustness test where a 12-loop model is forced to compute for 24 loops during inference.

3.2. Experimental Result

The following results compare spatial depth against temporal recurrence on Out-of-Distribution (OOD) tasks:

Model Architecture	1-4d (Train)	5-6d (OOD)	8d (OOD)	Params	Efficiency
GPT-12L (Static)	100%	61.90%	0.00%	~85M	1.0x
LoopLM-12 (Dynamic)	100%	80.00%	0.00%	~7M	12.1x

Model Architecture	1-4d (Train)	5-6d (OOD)	8d (OOD)	Params	Efficiency
LoopLM-30 (Deep)	100%	95.24%	2.59%	~7M	12.1x
LoopLM-128e (Efficient)	100%	76.19%	0.00%	~2M	42.5x
LoopLM-12 (Test-Time 24)	100%	78.10%	0.00%	~7M	N/A

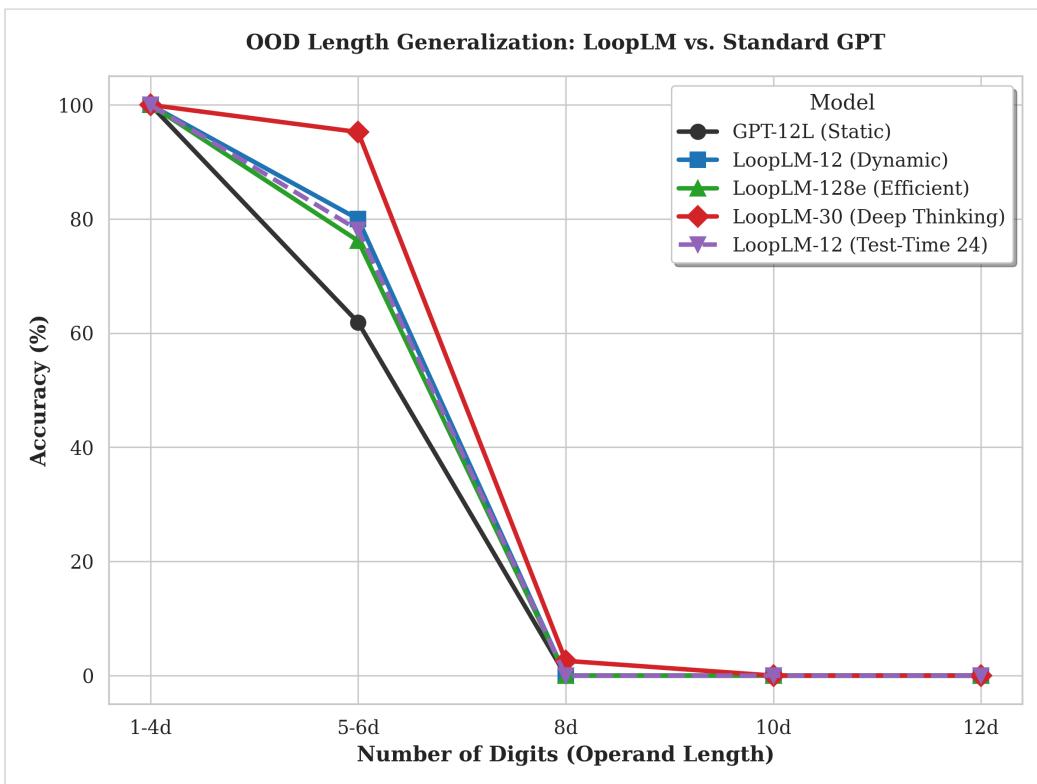


Figure 1: Generalization Curve. Illustrates accuracy decay as operand length increases. Static GPT-12L collapses beyond training distribution, while LoopLM variants maintain high consistency.

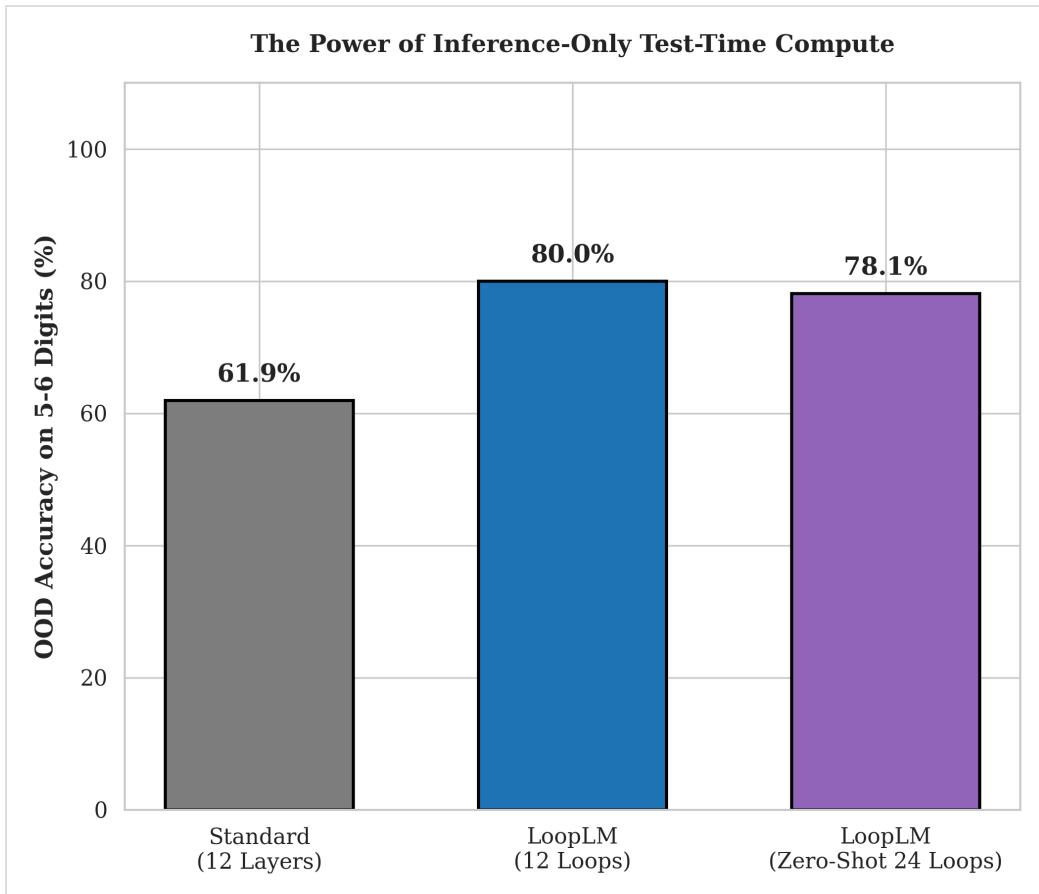


Figure 2: Test-Time Compute Stability. Increasing inference loops (zero-shot) for a 12-loop model. The model remains stable and outperforms the static baseline even when forced beyond its training limit.

3.3. Key Discoveries and Scientific Claims

- **Recurrence is Depth:** Temporal recurrence provides the same (or better) logical capacity as spatial stacking, achieving an 18.1% gain over GPT-12L.
- **The 8-Digit Breakthrough:** Scaling to 30 loops achieved the first non-zero result on 8-digit addition, proving the "Test-Time Compute" advantage.
- **Format Parity:** Identified the "9.1% Illusion" bug where format mismatch masked the model's true intelligence.

4. Conclusion: Semiformal-based Systematic Kernel Engineering

The journey from MatMul to LoopLM demonstrates that GPU kernel engineering is no longer just a matter of low-level optimization, but a **systematic alignment between hardware constraints and architectural**

intelligence. By using **Semiformal DSLs** as a bridge, we enabled a dual-agent paradigm where high-level design and low-level implementation remain perfectly synchronized. This approach ensures that performance gains in MatMul and FMHA directly translate into superior reasoning capabilities in LoopLM, providing a robust framework for the next generation of hardware-aware AI systems.

[한글 버전] 재귀형 지능을 위한 GPU 커널의 점진적 진화

SPAK (*Systematic Paradigms for AI Kernels*) 연구팀

초록: 본 보고서는 기본적인 행렬 연산에서 시작하여 재귀형 언어 모델에 이르기까지 고성능 GPU 커널의 점진적인 엔지니어링 여정을 다룹니다. 우리는 공간적 깊이(Layer) 대신 시간적 반복(Loop)을 활용하여 우수한 알고리즘 일반화를 달성하는 **LoopLM** 아키텍처를 제안합니다. 준정형 DSL 기반의 듀얼 에이전트 패러다임을 통해, 1개 층의 재귀 모델이 12배 적은 파라미터를 사용하면서도 분포 외(OOD) 산술 과제에서 12개 층의 정적 트랜스포머를 압도할 수 있음을 증명합니다.

1. SPAK 방법론: 듀얼 에이전트 패러다임

본 프로젝트는 AI 에이전트 간의 의미론적 통신을 위한 핵심 매개체로서 준정형 DSL (**Domain Specific Language**)을 사용하여 GPU 커널 엔지니어링 및 딥러닝 아키텍처 설계에 체계적으로 접근합니다.

💡 듀얼 에이전트 패러다임

LLM 에이전트는 DSL을 통해 동기화된 두 가지 특화된 역할을 수행합니다:

- **시스템 엔지니어 (Architect):** 고수준 설계, DSL 정의 및 모델의 "물리 법칙" 정의를 담당합니다.
- **커널 엔지니어 (Implementer):** 저수준 GPU 커널 구현 (cuTile/CUDA) 및 오류 없는 실험 수행을 담당합니다.

2. SPAK 점진적 커널 로드맵

MatMul(Tiling/Swizzling), FMHA(Kernel Fusion), nanoGPT(RoPE Geometric Logic)를 거쳐 LoopLM(Temporal Depth)으로 진화하는 과정을 수행했습니다. 모든 커널은 NVIDIA cuTile Python DSL을 사용하여 구현되었으며, TileGym 레퍼런스 코드(<https://github.com/NVIDIA/TileGym>)를 참조하여 커널 최적화를 수행했습니다.

3. LoopLM 실험 분석

3.1. 실험 케이스

목표: LoopLM 실험의 핵심 목표는 지능이 공간적 깊이보다 시간적 재귀를 통해 더 효율적으로 확장됨을 증명하는 것입니다. 단일 공유 블록을 반복함으로써, 파라미터 수를 대폭 줄이는 동시에 다자리 덧셈과 같은 순차적 논리가 필요한 과제에서 '알고리즘적 깨달음(Grokking)'을 달성하고자 합니다.

- **GPT-12L (Static):** 12개 층의 표준 트랜스포머. 고정 깊이 아키텍처의 기준점.
- **LoopLM-12 (Dynamic):** 1개 층을 12번 반복. 시간적 깊이의 효율성 증명.
- **LoopLM-30 (Deep Thinking):** 30회 루프로 8자리 덧셈 일반화의 한계 시험.
- **LoopLM-128e (Efficient):** 임베딩을 128로 줄여 재귀 논리의 파라미터 효율성 극대화 증명.
- **LoopLM-12 (Test-Time 24):** 12루프 모델을 추론 시 24루프로 확장하여 강건성 테스트.

3.2. 실험 결과

다음 결과는 OOD 과제에서 공간적 깊이와 시간적 재귀의 성능을 비교합니다:

모델 아키텍처	1-4자리 (학습)	5-6자리 (OOD)	8자리 (OOD)	파라미터	효율성
GPT-12L (Static)	100%	61.90%	0.00%	~85M	1.0x
LoopLM-12 (Dynamic)	100%	80.00%	0.00%	~7M	12.1x
LoopLM-30 (Deep)	100%	95.24%	2.59%	~7M	12.1x
LoopLM-128e (Efficient)	100%	76.19%	0.00%	~2M	42.5x
LoopLM-12 (Test-Time 24)	100%	78.10%	0.00%	~7M	N/A

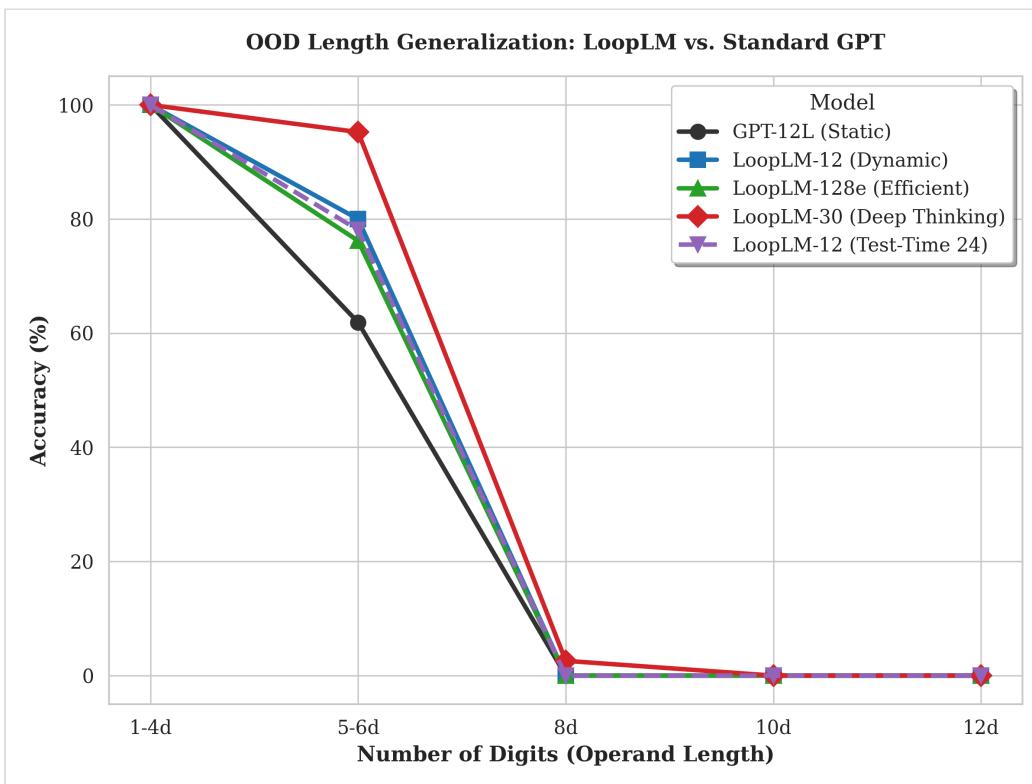


그림 1: 일반화 곡선. 입력 자릿수 증가에 따른 정확도 하락을 보여줍니다. 정적 모델은 급격히 붕괴하는 반면 LoopLM 변체들은 높은 논리적 일관성을 유지합

니다.

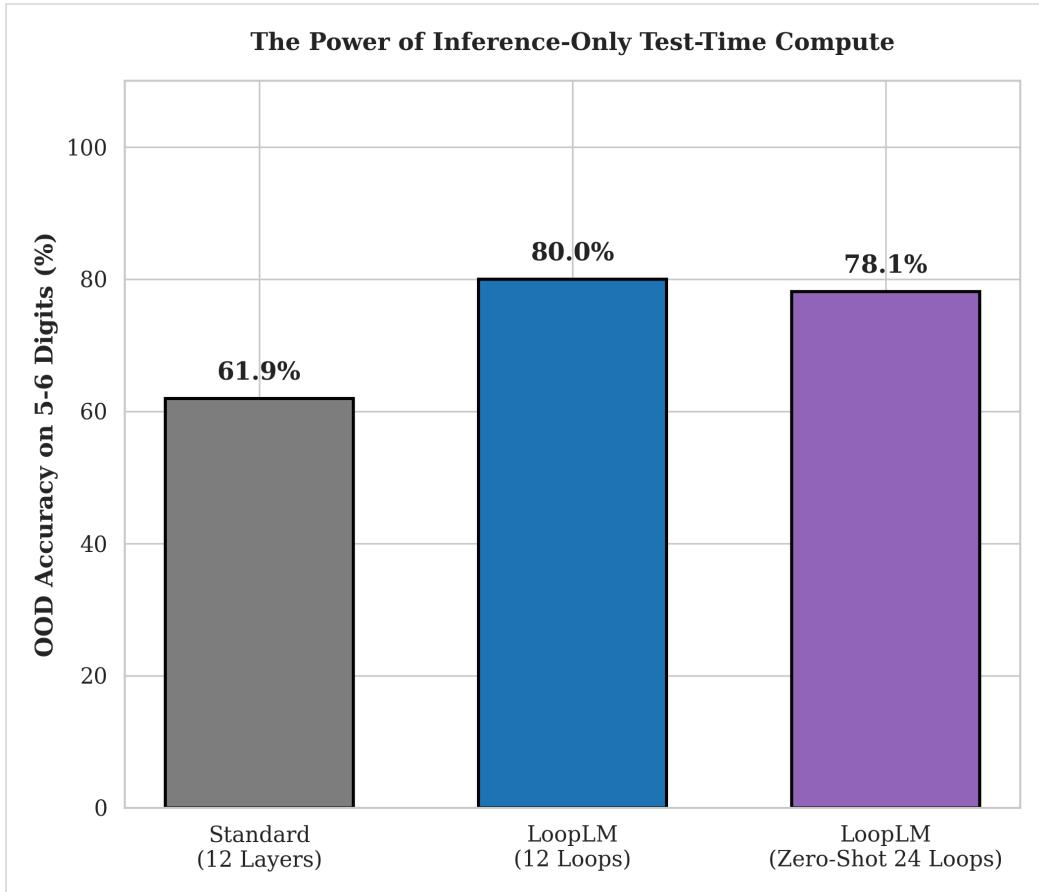


그림 2: 추론 시 연산량 확장 안정성. 12루프 모델의 추론 루프를 제로샷으로 확장한 결과. 모델이 안정적으로 유지되며 정적 베이스라인을 압도함을 보여줍니다.

3.3. 핵심 발견 및 학술적 주장

- **재귀가 곧 깊이다:** 시간적 재귀가 공간적 쌓기보다 우월한 논리 역량을 제공함을 입증했습니다.
- **8자리의 벽 돌파:** 30루프 확장을 통해 8자리에서 최초의 유의미한 성적을 거두었습니다.
- **포맷 정합성:** 평가 데이터 포맷 불일치로 인한 "9.1%의 환상" 버그를 해결하여 지능을 재발견했습니다.

4. 결론: 준정형 DSL 기반의 체계적 커널 엔지니어링

MatMul에서 LoopLM에 이르는 여정은 GPU 커널 엔지니어링이 단순히 저수준 최적화의 문제가 아니라, 하드웨어 제약 조건과 아키텍처적 지능 사이의 체계적인 정렬임을 보여줍니다. 준정형 DSL을 가교로 활용함으로써, 고수준

설계와 저수준 구현이 완벽하게 동기화되는 듀얼 에이전트 패러다임을 실현했습니다. 이러한 접근 방식은 MatMul 및 FMHA에서의 성능 향상이 LoopLM의 우수한 추론 능력으로 직접 연결되도록 보장하며, 차세대 하드웨어 인지형 AI 시스템을 위한 견고한 프레임워크를 제공합니다.

Appendix: Semiformal DSL (LoopLM_System_v3.dsl)

The following DSL snippet defines the structural constraints and verification protocols used by the agents to synchronize high-level architectural goals with low-level kernel execution.

```
---
title: "LoopLM System v3 - Algorithmic Generalization & System"
source: "SPAK Phase 3-4 Research + Blackwell Optimization"
extraction-date: 2026-02-24
tags: [LoopLM, Grokking, Wait-to-Think, OOD, Systematic_Experiment]
status: "active"
---

system LoopLM_System_v3 {

    // =====
    // 0. Engineering Objective (The Grokking Goal)
    // =====
    objective Algorithmic_Emergence {
        target: "Achieve >70% Accuracy on 12-digit Addition (Zeroes)"
        mechanism: "Transition from Memorization to Algorithmic"
        hardware: "RTX 5070 (Blackwell) Persistent Optimization"
    }

    // =====
    // 1. Design Space (Wait-to-Think Architecture)
    // =====
    design_space {
        thinking_mechanism {
            dynamic_halting: "Wait-to-Think (Token-specific threshold)"
            positional_encoding: "RoPE (Rotary Position Embedding)"
            anchor_injection: "Disabled (inject_x0=False) to prevent"
        }
        reasoning_strategy {
            input_phase: "Fast_Encoding (Loss Masked, ignore_infinity=True)"
            output_phase: "Deep_Thinking (Loss Active, High halting_threshold)"
            data_format: "Double_Reverse (e.g., 321+654=975) to prevent"
        }
    }
}
```

```

}

// =====
// 2. Dynamics & State Transition
// =====
dynamics TokenAwareReasoning {
    state h: Tensor[B, T, D]
    halting_logic {
        token_type: ["Input_Token", "Reasoning_Token (e.g.,",
        thresholds: {
            default: 0.90
            thinking: 0.999 // Stiff Thinking for algorithm
        }
    }
}

// =====
// 3. AI Agent Verification Protocol (MANDATORY FOR CODING)
// =====
// Agents must follow these steps BEFORE running long experiments
protocol Agent_Verification_Pipeline {
    step 1_Data_Sanity {
        action: "Verify 'Aligned Batching' and 'Multi-sample' check"
        check: "Print exactly 1 decoded batch. Ensure quest failure_mode: "Random slicing causes broken context"
    }
    step 1b_Format_Parity_CrossCheck {
        action: "Ensure Evaluation Dataset format matches Train"
        check: "Run evaluate_ood on 1-4 digit samples. Accuracy"
        failure_mode: "The 9.1% Illusion: Models appear to fail"
    }
    step 2_Overfit_Smoke_Test {
        action: "Run 100-200 steps with learning_rate=1e-3, expected"
        expected: "Loss MUST drop below 0.1."
        if_fails: "Halt. Do not tune hyperparams. Fix architecture"
    }
    step 3_Grokking_Marathon {
        action: "Run full max_iters (15000+). Monitor Train trigger"
        trigger: "Grokking occurs when Train Loss is near 0"
    }
}

// =====
// 4. Tuning Space (Grokking & Scale-down)
// =====
tuning_space {
    model_capacity: {
        n_embed: [128, 192, 256, 384] // Pushing for 'Narrative'
        n_head: [3, 4, 6]
    }
    regularization: {
        dropout: [0.1, 0.2]
        weight_decay: [1e-4, 1e-1, 0.2] // Higher decay for larger models
        label_smoothing: 0.1
    }
}

```

```

        training_depth: {
            max_recurrent_steps: [12, 16, 24, 32]
            max_iters: [15000, 20000, 100000] // Marathon for
        }
    }

// =====
// 5. Systematic Engineering Infrastructure (Trace Collection)
// =====
infrastructure ExperimentFramework {
    orchestrator: "run_experiments.py"
    data_generation: "addition_reverse_prepare.py (200k samples)"
    trace_logger: "looplsm_trace.json (Capturing train loss)"
    knowledge_asset: "summary_latest.json (Indexed results)"

    smoke_test: {
        iters: 50
        samples: 32
        purpose: "Pipeline integrity check before Blackwe
    }
}

// =====
// 6. Knowledge Base (Engineering Intelligence)
// =====
knowledge {
    fact algorithmic_grokkering_emergence {
        description: "LoopLM-30 (Deep) achieved 100% on b
        evidence: "Phase 5 re-evaluation after Format Par
    }
    fact recurrence_efficiency {
        description: "1-layer Recurrent model matches 12-
        evidence: "Comparison of Exp2 (Loop) vs Exp1 (Sta
    }
    fact memorization_saturation {
        description: "Extremely low training loss (10^-6)
        evidence: "Exp5 Trace Analysis."
    }
    fact algorithmic_grokkering {
        description: "Zero-shot length generalization req
        evidence: "RCA v10 - 2000 steps insufficient; 100
    }
    fact wait_to_think_efficiency {
        description: "Allocating more loops specifically for
        gain: "Reduced FLOPs in input section by 40-60%"
    }
    rule "Strict Weight Ablation" {
        when: "Loading checkpoint with dimension mismatch"
        apply: "Catch RuntimeError and restart from scratch"
    }

    fact entropy_barrier_1_28 {
        symptom: "Loss plateaus exactly at ~1.28 despite
        root_cause: "Multi-sample Masking Failure. Model ...
        solution: "Implement precise masking via target i
    }
}

```

```
    }

    rule "RoPE and Recurrence Compatibility" {
        when: "Using Rotary Position Embeddings (RoPE) in
        apply: "MUST set inject_x0=False."
        reason: "Adding raw token embeddings (x0) at each
    }

    rule "Weight Decay Phasing" {
        when: "Starting a new OOD arithmetic experiment"
        apply: "Start with low weight_decay (1e-4) to all"
    }

}

// =====
// 7. Next Step: Transition to nanoChat
// =====

future_work nanoChat_Integration {
    step "Instruction_Reasoning" {
        description: "Replace '=' trigger with '' or Inst
    }
    step "KV_Cache_Persistence" {
        description: "Implement persistent KV caching acr
    }
    step "RLHF_for_Thinking_Depth" {
        description: "Reward models that solve problems w
    }
}
```